


Linear Regression with Gradient Descent

Linear Regression Model



- $f(x) = w_0 + w_1x$
- How to choose the optimal weight vector ?
 - Iterative method (Gradient Descent)
 - Closed form solution (OLS)

Gradient Descent Algorithm

Repeat until convergence:

$$w_j = w_j - \eta \frac{\partial J(w)}{\partial w_j} \quad j = 0, 1, \dots, d$$

}

- η = learning rate
- $J(w)$ = cost function = $\frac{1}{2m} \sum_{i=1}^m (f(x^{(i)}) - y_i)^2$
- $\frac{\partial J(w)}{\partial w_j} = \frac{1}{m} \sum_{i=1}^m [f(x^{(i)}) - y_i] x_j^{(i)}$
 - d = number of dimensions
 - m = number of training examples
 - $x^{(i)}$ = input variable (features) of the i^{th} training example
 - y_i = output variable of the i^{th} training example
 - $x_j^{(i)}$ = the j^{th} feature in the i^{th} training example

Multivariate Linear Regression

- Linear Regression with multiple variables (features)

Size	No bedrooms	No bathrooms	Price (\$1000)
1400	3	2	250
1719	3	2	299
...

- Feature vector and weight vector have $d+1$ dimensions
- $f(\mathbf{x}) = \mathbf{w}^T \mathbf{x} = w_0x_0 + w_1x_1 + \dots + w_dx_d$
- Cost function = $J(\mathbf{w}) = \frac{1}{2m} \sum_{i=1}^m (f(\mathbf{x}^{(i)}) - y_i)^2$

Repeat until convergence:

$$w_j = w_j - \eta \frac{\partial J(\mathbf{w})}{\partial w_j} \quad j = 0, \dots, d$$

}

$$\text{where } \frac{\partial J(\mathbf{w})}{\partial w_j} = \frac{1}{m} \sum_{i=1}^m [f(\mathbf{x}^{(i)}) - y_i] x_j^{(i)}$$

Stochastic Gradient Descent

- Batch gradient descent uses all training examples for each iteration.
 - Slow, requires large memory but less sensitive to variation
- Stochastic gradient descent updates the parameters with every training example
 - Much faster, allows on-line training but may suffer from high variance

shuffle data set

repeat {

 for $i = 1, \dots, m$ {

$$w_j = w_j - \eta \nabla_{w_j} J \text{ where } \nabla_{w_j} J = [f(\mathbf{x}^{(i)}) - y_i] x_j^{(i)}$$

 ($j = 0, 1, \dots, d$)

 }

}

Mini Batch Gradient Descent

- Similar to stochastic gradient descent, but the parameters are updated with n training examples for each iteration
 - Batch Gradient Descent: m training examples per iteration
 - Stochastic Gradient Descent: 1 training example per iteration
 - Mini Batch Gradient Descent: n training examples per iteration, $n \ll m$