


Logistic Regression

Logistic regression in one dimension

a) Example: APACHE II Score and Mortality in Sepsis

The following figure shows 30 day mortality in a sample of septic patients as a function of their baseline APACHE II Score. Patients are coded as 1 or 0 depending on whether they are dead or alive in 30 days, respectively.



Sigmoid Function

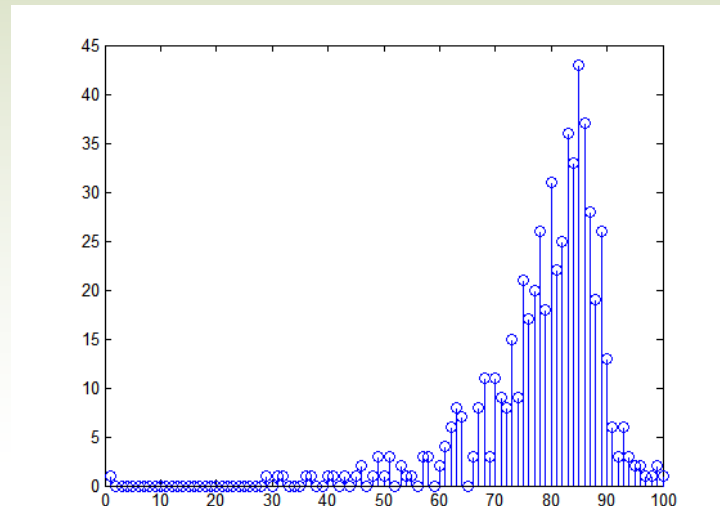
- Bernoulli distribution has 2 possible outcomes

- PMF: $p(x) = y^x(1 - y)^{1-x}$

- Recall for linear regression

$$y(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + \varepsilon \text{ where } \varepsilon \sim \text{Gaussian}$$


- Error Histogram



Sigmoid Function

- Bernoulli distribution has 2 possible outcomes
 - PMF: $p(x) = y^x(1 - y)^{1-x}$
- If the response is binary, $y \in \{0,1\}$, then $\varepsilon \sim \text{Bernoulli}$
 $p(y|\mathbf{x}, \mathbf{w}) = \text{Ber}(y|\mathbf{w}^T \mathbf{x})$
- Pass $\mathbf{w}^T \mathbf{x}$ through function $f(\mathbf{w}^T \mathbf{x})$ such that $0 \leq f(\mathbf{w}^T \mathbf{x}) \leq 1$
- For logistic regression, choose sigmoid (logistic) function

$$\text{sigm}(\mathbf{w}^T \mathbf{x}) = \frac{1}{1 + \exp(-\mathbf{w}^T \mathbf{x})}$$



Weight Vector

- How to choose optimal weight vector \mathbf{w} ?
 - Gradient descent (ascent) method
 - $\mathbf{w} = \operatorname{argmin}_{\mathbf{w}} J(\mathbf{w}) \rightarrow \mathbf{w} = \mathbf{w} - \eta \frac{\partial J(\mathbf{w})}{\partial \mathbf{w}}$ if $J(\mathbf{w})$ is convex
 - $\mathbf{w} = \operatorname{argmax}_{\mathbf{w}} J(\mathbf{w}) \rightarrow \mathbf{w} = \mathbf{w} + \eta \frac{\partial J(\mathbf{w})}{\partial \mathbf{w}}$ if $J(\mathbf{w})$ is not convex
- With 1 training example, formulate the cost function
 - If model predicts correctly, “reward” the model, cost = 0
 - Else “penalize” heavily, cost = ∞
 - Cost = $-[y \log(f(x)) + (1 - y) \log(1 - f(x))]$
- With N training examples
$$J = -\sum_{i=1}^N y^{(i)} \log(f(\mathbf{x}^{(i)})) + (1 - y^{(i)}) \log(1 - f(\mathbf{x}^{(i)}))$$
- Is J convex ?

Stochastic Gradient Descent

shuffle data set randomly

repeat {

 for $i = 1, \dots, N$ {

$$w_j = w_j - \eta \nabla_{w_j} J \text{ where } \nabla_{w_j} J = [\text{sigm}(\mathbf{w}^T \mathbf{x}^{(i)}) - y^{(i)}] x_j^{(i)}$$

 ($j = 0, 1, \dots, d$)

 }

}

- $x_j^{(i)}$: j -th feature of the i -th training example

- Recall the update equation for Linear Regression

$$w_j = w_j - \eta \nabla_{w_j} J \text{ where } \nabla_{w_j} J = [f(\mathbf{x}^{(i)}) - y^{(i)}] x_j^{(i)}$$

- Once the optimal weight vector is obtained, use the sigmoid function to classify a test data


The bias

- The bias plays a role in shifting the sigmoid curve

- $s = \text{sigm}(z) = \frac{1}{1 + \exp(-z)}$

Where $z = \mathbf{w}^T \mathbf{x} = w_0 x_0 + w_1 x_1 + w_2 x_2 + \dots$

Note: $x_0 = 1$



Logistic Regression Classification Example with 2-D Dataset

- Assume $\mathbf{w}^T = [w_0 \quad w_1 \quad w_2] = [0.5 \quad -1.3 \quad 3.2]$
- If $s > 0.5 \rightarrow$ Class 1, else \rightarrow Class 0
 - $s = \text{sigm}(z) = \frac{1}{1+\exp(-z)}$
 $z = \mathbf{w}^T \mathbf{x} = w_0 x_0 + w_1 x_1 + w_2 x_2$
 $x_0 = 1$

x_1	x_2	y	z	s	\hat{y}
4.1	1.3	0	-0.67	0.34	0
4.5	1.5	1	-0.55	0.38	0*
1.7	0.4	0	-0.43	0.39	0
0.5	0.7	1	2.09	0.89	1

* Misclassified

Softmax Regression

- Generalization of logistic regression for multi-class classification
- a.k.a. Multinomial Logistic Regression, Maximum Entropy Classifier
- Logistic regression

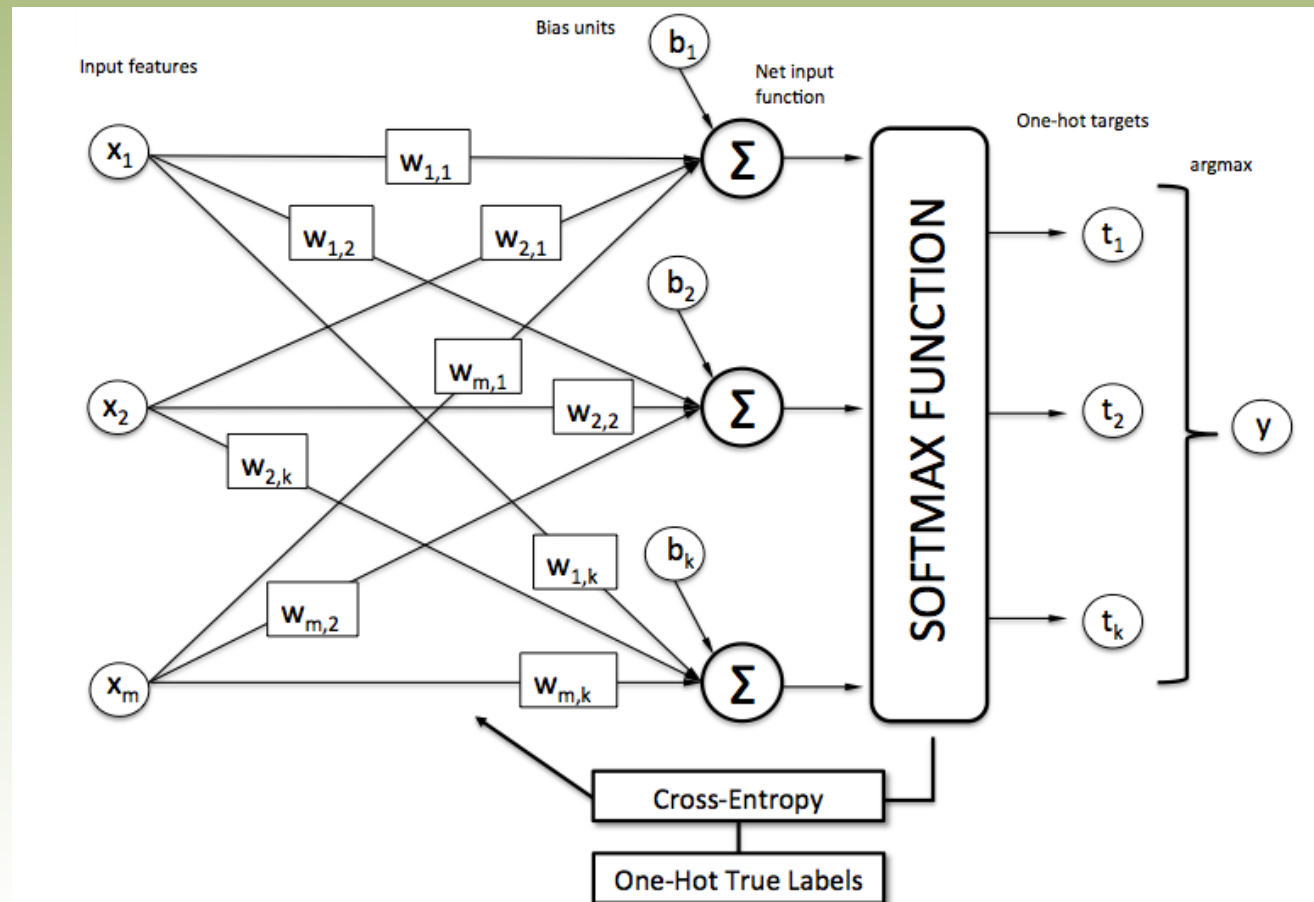
$$s = \text{sigm}(z) = \frac{1}{1 + \exp(-z)} = \frac{\exp(z)}{1 + \exp(z)}$$

Where $z = \mathbf{w}^T \mathbf{x}$

- Softmax regression

$$\phi = \text{softmax}(\mathbf{z}) = \frac{\exp(z_j)}{\sum_{i=1}^K \exp(z_i)} = p(y = j | \mathbf{z})$$

$p(y = j | \mathbf{z}) \rightarrow$ probability of class j



Softmax Regression

Reference by S. Raschka

Training

- Cross entropy

$$H(p, q) = - \sum_i p_i \log(q_i)$$

- Update equation for each class

$$\mathbf{w}_j = \mathbf{w}_j - \eta \nabla_{\mathbf{w}_j} J$$

$$J = -\frac{1}{N} \sum_{i=0}^N H(y_i, \phi(\mathbf{x}^{(i)}))$$

$$\nabla_{\mathbf{w}_j} J = \frac{1}{N} \sum_{i=0}^N [\{\phi(\mathbf{x}^{(i)}) - y^{(i)}\} x_j^{(i)}] \quad j \in \{1, \dots, D\}$$

- $y^{(i)}$: true class label
- $\phi(\mathbf{x}^{(i)})$: softmax output (not the predicted label)
- J : average of all cross entropies over all training examples
- \mathbf{w}_j : weight vector for j^{th} feature
- $x_j^{(i)}$: j^{th} feature in i^{th} training example

Prediction

- Compute a score for each class

$$s_j(\mathbf{x}) = p(y = j | \mathbf{z}) = \frac{\exp(z_j)}{\sum_{i=1}^K \exp(z_i)}$$

Where $\mathbf{z} = \mathbf{w}^T \mathbf{x}$

- Choose the class with the highest score

$$\hat{y} = \operatorname{argmax}_j s_j(\mathbf{x})$$